

John Boaz Lee

jtlee@wpi.edu | johnboaz.github.io

Senior machine learning engineer and researcher working primarily on deep learning, sequence and graph modeling, representation learning, and data mining.

Education

08/2015—12/2019 **Ph.D., Worcester Polytechnic Institute**

Computer Science, GPA 4.0/4.0

- *Thesis:* Deep Learning on Graph-structured Data
- *Committee:* Xiangnan Kong (advisor), Elke Rundensteiner, Yanhua Li, Ryan Rossi

06/2010—06/2012 **M.Sc., University of the Philippines – Diliman**

Computer Science, GPA 4.0/4.0

- *Thesis:* A Link Prediction Algorithm for Heterogeneous Bibliographic Information Networks
- *Committee:* Henry Adorna (advisor), Prospero Naval Jr., Jaderick Pabico, Rommel Feria

06/2004—06/2008 **B.Sc., Silliman University**

Computer Science, GPA 3.1/4.0

- *Thesis:* Dumaguete e-Traveler – A Knowledge-based Decision Support System
- *Advisor:* Dave Marcial

Experience

01/2020—Present **Facebook**

Senior Machine Learning Engineer, Menlo Park

- Recently working primarily on graph learning for ads retrieval and ranking. I have experience with a wide range of graph-based methods, and these include GNNs (\$31.5M/year in revenue), non-parametric algorithms like node2vec and Personalized PageRank (\$370M/year in revenue), and low-latency random walk methods that can serve real-time use cases (\$120M/year in revenue).
- Led the engineering effort to deploy user intent prediction models across 3 surfaces (jobs, marketplace, shops). These were adopted in 13 separate use cases, most notably for promotions targeting (\$2.4M/year in savings) and ads ranking (\$2.8M/year in revenue). The models were also found to be 24% better than existing methods.
- Launched >30 ML solutions for a wide variety of use cases, including: graph representation learning for product recommendation (+5.7% transactions), multimodal topic prediction (+3% P@R85), and sequence modeling for anti scraping (2x recall at fixed precision; multi-task, multi-label variant reduced training and deployment time by 80%; helped take down 2 malicious third-party apps with >100M downloads).
- Core code contributor to an internal PyTorch framework for sequence modeling.
- Main responsibilities included building and maintaining data pipelines, performing adhoc data analyses, and building, testing, and deploying machine learning models.

06/2019—08/2019 **Facebook**

Machine Learning Researcher, Menlo Park (Internship)

- Deployed a recurrent pipeline which generated person-level graph embeddings for >1B users — leveraging their interactions from across FB's family of apps.
- Demonstrated usefulness of approach by showing consistent AUROC gains across 9 tasks.

- 05/2018—08/2018 **Adobe**
Machine Learning Researcher, San Jose (Internship)
 - Worked on a deep learning model for entity alignment which was deployed internally. This resulted in two research publications and a patent (approved in 2023).
- 05/2017—08/2017 **Xerox PARC**
Machine Learning Researcher, Palo Alto (Internship)
 - Led a project on graph classification and collaborated on a second one on dynamic network representation learning. Published two research papers from these projects.
- 08/2015—06/2019 **Worcester Polytechnic Institute**
Research Assistant, Worcester
 - Researched and developed deep learning techniques for various types of relational data. Results were published at top data mining venues like TKDE, TKDD, KDD, and CIKM.
 - Assisted professors in the following courses: Data Mining, Computer Networks, Software Security, Temporal Logic, and Theory of Computation (teaching assistant during initial two years at WPI).
- 10/2012—10/2012 **Nara Institute of Science and Technology**
Research Assistant, Nara (Internship)
 - Proposed and built a system for automatic software patch reviewer recommendation.
 - Finished a draft research publication within a month, the paper was later accepted for publication and won a “Best Presentation Paper” award.
- 06/2012—07/2015 **Ateneo de Manila University**
Instructor, Metro Manila
 - Taught the following college-level courses: Data Structures and Algorithms, Java Programming, C# Programming, Design Patterns, and Linux Systems Programming.
 - Coached, in a minor capacity, competitive programming team (13/53 in ICPC regionals).
- 06/2010—06/2012 **University of the Philippines – Diliman**
Instructor, Metro Manila
 - Handled the following courses: C Programming, Software Engineering, Data Structures and Algorithms. Taught primarily to an audience without a computing background.
- 04/2008—04/2009 **Enterprise Solutions**
Software Engineer, Cebu
 - Designed and developed core tools used in the application development framework of an ERP system. Promoted after six months at the company.
- 04/2006—06/2006 **Neri & Hu Design and Research Office**
Software Engineer, Shanghai (Internship)
 - Designed and implemented an in-house Inventory Management System from scratch.

Publications

Citation statistics available on [Google Scholar](#)

Published at top data mining venues including: KDD (3), SDM (3), TKDD (2), TKDE (1), and CIKM (1)

Patents: [US 11,544,535 B2](#)

TKDE'22 Role-based Graph Embeddings

N. K. Ahmed, R. A. Rossi, **J. B. Lee**, T. L. Willke, R. Zhou, X. Kong, H. Eldardiry
IEEE Transactions on Knowledge and Data Engineering. 2022.

KDD'21 Energy-efficient Models for High-dimensional Spike Train Classification using Sparse Spiking Neural Network

H. Yin, **J. B. Lee**, X. Kong, T. Hartvigsen, S. Xie
ACM SIGKDD International Conference on Knowledge Discovery and Data Mining. 2021.

- KDD'20** Recurrent Networks for Guided Multi-Attention Classification
X. Dai, X. Kong, T. Guo, **J. B. Lee**, X. Liu, C. Moore
ACM SIGKDD International Conference on Knowledge Discovery and Data Mining. 2020.
- TKDD'20** On Proximity and Structural Role-based Embeddings in Networks: Misconceptions, Techniques, and Applications
R. A. Rossi, D. Jin, S. Kim, N. K. Ahmed, D. Koutra, **J. B. Lee**
ACM Transactions on Knowledge Discovery from Data. 2020.
- SDM'20a** Deep Parametric Model for Discovering Group-cohesive Functional Brain Regions
J. B. Lee, X. Kong, C. M. Moore, N. K. Ahmed
SIAM International Conference on Data Mining. 2020.
- SDM'20b** Dual-Attention Recurrent Networks for Affine Registration of Neuroimaging Data
X. Dai, X. Kong, X. Liu, **J. B. Lee**, C. Moore
SIAM International Conference on Data Mining. 2020.
- TKDD'19** Attention Models in Graphs: A Survey
J. B. Lee, R. A. Rossi, S. Kim, N. K. Ahmed, E. Koh
ACM Transactions on Knowledge Discovery from Data. 2019.
- CIKM'19** Graph Convolutional Networks with Motif-based Attention
J. B. Lee, R. A. Rossi, X. Kong, S. Kim, E. Koh, A. Rao
ACM International Conference on Information and Knowledge Management. 2019.
- ANS'19** Learning Compact Graph Representations via an Encoder-Decoder Network
J. B. Lee, X. Kong
Applied Network Science. 2019.
- BigData'18** Dynamic Network Embeddings: From Random Walks to Temporal Random Walks
G. Nguyen, **J. B. Lee**, R. Rossi, N. Ahmed, E. Koh, S. Kim
IEEE International Conference on Big Data. 2018.
- KDD'18** Graph Classification using Structural Attention
J. B. Lee, R. A. Rossi, X. Kong
ACM SIGKDD International Conference on Knowledge Discovery and Data Mining. 2018.
- BigNet @ WWW'18** Continuous-Time Dynamic Network Embedding
G. Nguyen, **J. B. Lee**, R. Rossi, N. Ahmed, E. Koh, S. Kim
International Workshop on Learning Representations for Big Networks @ The Web Conference. 2018.
- SDM'17** Identifying Non-linear Contrasting Networks from Time-Series Data: Application to Brain Network Analysis
J. B. Lee, X. Kong, Y. Bao, C. Moore
SIAM International Conference on Data Mining. 2017.
- IWESEP @ APSEC'13** Patch Reviewer Recommendation in OSS Projects
J. B. Lee, A. Ihara, A. Monden, K. Matsumoto
International Workshop on Empirical Software Engineering in Practice @ Asia-Pacific Software Engineering Conference. 2013. (Best Presentation Award)
- JoC'13** Understanding the Behavior of Filipino Twitter Users during Disaster
J. B. Lee, M. Ybañez, M. M. de Leon, M. R. E. Estuar
Journal on Computing. 2013.
- CGAT'13** Characterizing Behavior and Features of Participants and Observers during Disaster on Twitter
J. B. Lee, M. Ybañez
International Conference on Computer Games Multimedia and Allied Technologies. 2013.

- ASONAM'12 Link Prediction in a Modified Heterogeneous Bibliographic Network
J. B. Lee, H. Adorna
International Conference on Advances in Social Networks Analysis and Mining. 2012.
- JUCS'12 Uncovering the Social Dynamics of Online Elections
J. B. Lee, G. Cabunducan, R. Castillo, F. G. Cabarle, J. A. Malinao
Journal of Universal Computer Science. 2012.
- JIP'12 A Quantitative Analysis-based Algorithm for Optimal Data Signature Construction of Traffic Data Sets
 J. A. Malinao, R. A. B. Juayong, E. R. F. Oquendo, R. M. U. Tadas, **J. B. Lee**, J. B. Clemente, M. S. Gaabucayan-Napalang, J. R. F. Regidor, H. N. Adorna
Journal of Information Processing. 2012.
- ASONAM'11 Voting Behavior Analysis in the Election of Wikipedia Admins
 G. Cabunducan, R. Castillo, **J. B. Lee**
International Conference on Advances in Social Networks Analysis and Mining. 2011. (Best Poster Paper)

Professional Service

- Program Committee* **WWW**: The Web Conference (2022)
CIKM: International Conference on Information and Knowledge Management (2020)
BigData: IEEE International Conference on Big Data (2018, 2020)
- Journal Reviewer* **TKDD**: ACM Transactions on Knowledge Discovery from Data (2017, 2021)
TNNLS: IEEE Transactions on Neural Networks and Learning Systems (2021)
VLDB: International Journal on Very Large Data Bases (2020)
TKDE: IEEE Transactions on Knowledge and Data Engineering (2017)
JUCS: International Journal of Universal Computer Science (2012)
- External Reviewer* **KDD**: Conference on Knowledge Discovery and Data Mining (2016 — 2019)
ICLR: International Conference on Learning Representations (2018 — 2019)
IJCAI: International Joint Conference on Artificial Intelligence (2017 — 2019)
ICDM: IEEE International Conference on Data Mining (2019)
SDM: SIAM International Conference on Data Mining (2017 — 2019)
NeurIPS: Conference on Neural Information Processing Systems (2016, 2018)
ASONAM: Conference on Adv. in Social Networks Analysis & Mining (2017 — 2018)
ICDE: IEEE International Conference on Data Engineering (2017)
AISTAT: International Conference on Artificial Intelligence and Statistics (2017)
CIKM: International Conference on Information and Knowledge Management (2016)
AAAI: AAAI Conference on Artificial Intelligence (2016)
- Other* Application Reviewer, *Facebook PhD Fellowship – Applied Statistics* (2021)
 Application Reviewer, *Facebook Request for Proposal – Applied Statistics* (2020)

Awards

- 2018 ***Graduate Research Innovation Exchange Awards Honorable Mention***, Worcester Polytechnic Institute CS Department
- 2013 ***Best Presentation Award***, IWESEP: International Workshop on Empirical Software Engineering in Practice (awarded to a single paper at workshop)
- 2012 ***Honors Graduate***, Phi Kappa Phi Honor Society (awarded to top 10% of graduating class)
- 2011 ***Best Poster Paper***, ASONAM: International Conference on Advances in Social Networks Analysis and Mining (awarded to a single poster paper)

Relevant Skills

Programming Python, Java, C#, C, SQL (Presto, Spark), LaTeX

Frameworks PyTorch, TensorFlow, Keras, Airflow, scikit-learn, NumPy, SciPy, pandas, matplotlib

Tools Jupyter Notebook, VS Code, GNU Screen, Vim, virtualenv, Git, Mercurial

Languages Cebuano (native), English (professional), Tagalog (professional), Mandarin (intermediate), Hokkien (intermediate)

References

A full list of references is available upon request

Updated March 23, 2024 ■